The form of dissipation shapes what a quantum reservoir remembers

Markus Baumann [1,2], Itamar Fink [2], Johannes Wittmann [2], Claudia Linnhoff-Popien [1,2], and Jonas Stein [1,2]

[1]Quantum Applications and Research Laboratory (QAR-Lab), Ludwig-Maximilians-Universität München, Munich, Germany
[2]Mobile and Distributed Systems Group, Ludwig-Maximilians-Universität München, Munich, Germany

Dissipation has emerged as a resource for quantum reservoir computing, with its rate as the usual control. Here we show that its form matters as much. It sets three properties of a memory: how long a stored signal survives, whether and how strongly new inputs can be written into it, and how strong a signal is left to measure. Independent decay shortens the memory as the loss grows. Collective decay has dark states in which the quantum Zeno effect holds the system, and stronger loss prolongs the memory. Whether an input reaches the dark states depends on the phases with which the qubits couple to the decay channel. Every new input also mixes with the inputs already stored. We prove that at a fixed input amplitude a stored signal combines, on average, as many inputs as the fraction of it that one input disturbs times the memory lifetime, so a longer memory tolerates only a weaker input. Conversely, we prove that when the amplitude falls as the inverse square root of the lifetime, single inputs keep a fixed share of the memory however long it lasts, and this limit reproduces our simulations without a fitted parameter. A weaker input leaves a fainter signal to measure, and at strong loss a collective channel whose phases differ from those of the input field recovers the most memory at every measurement budget. The form of the dissipation, the input amplitude and the measurement budget are joint design parameters of a physical memory.

Corresponding author
Markus Baumann markus.baumann@campus.lmu.de

1 Introduction

Processing information over time requires a memory of what came before. A physical system can provide such a memory because each input changes its state and leaves traces that influence its later response. Reservoir computing uses these traces to perform a task: an input stream drives a dynamical system, and a weighted sum of quantities describing its state produces the output. Only these weights are trained; this mapping is called a linear readout [1, 2]. In quantum reservoir computing (QRC), the evolving system is quantum mechanical, and the readout uses quantities obtained by measuring it [3]. The challenge is to keep earlier inputs accessible while new information continues to arrive.

The environment shapes how these traces fade. A system can lose energy or information to its surroundings through processes collectively described as dissipation [4–15]. Previous work has shown that tuning the dissipation rate improves memory, in some reservoirs with the best performance at an intermediate loss rate [6, 9]. Input, interactions, and dissipation can also be optimized together [16, 17]. The rate specifies how strongly the environment acts, but the way it acts is a separate design choice. Engineering this interaction can prepare and protect selected quantum states [18–21]. For a system processing an ongoing input stream, this raises the possibility of controlling which traces of the past survive.

This distinction becomes concrete for a reservoir of interacting spins, each a two-level quantum system. When the spins decay through separate routes into the environment, they undergo independent decay. When they share a common route, they undergo collective decay. In certain joint states, the different emission paths interfere and cancel, preventing decay through that shared route. These protected states are called

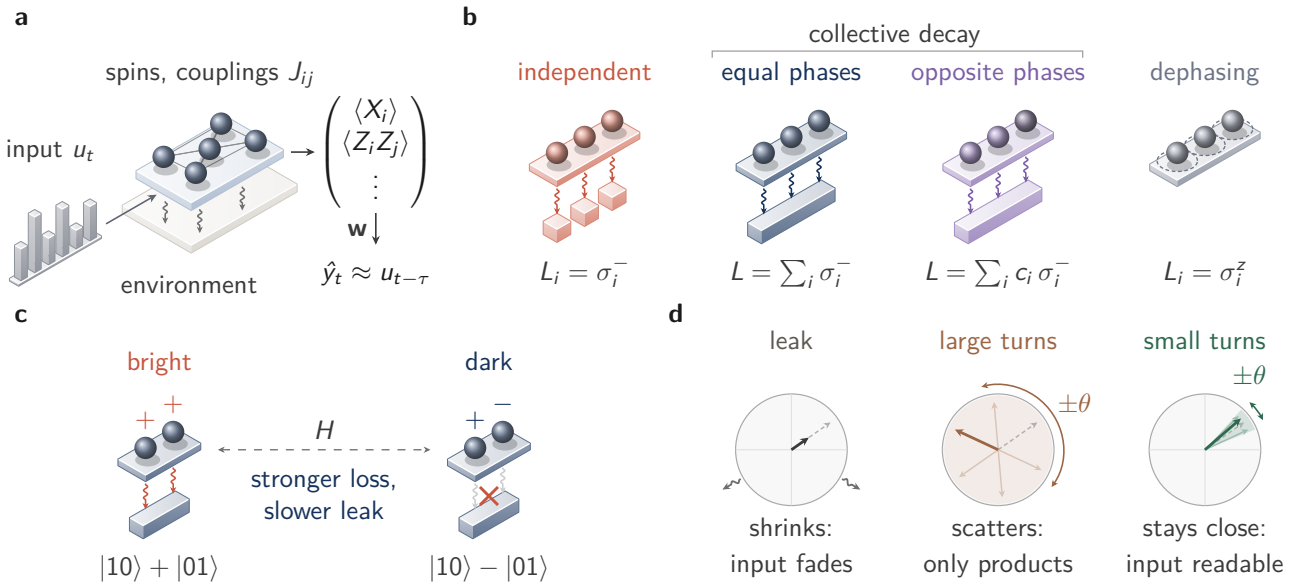


Figure 1: **Engineering how a quantum reservoir leaks. a** Inputs drive interacting spins, and a trained linear readout recalls past inputs. **b** Four forms of dissipation with the same total coupling. The darker spins couple with the opposite sign. **c** Two spins on one channel: their decays add (bright) or cancel (dark), and strong loss holds them in the dark state, the quantum Zeno effect. **d** A stored signal fades because it leaks away or because later inputs turn it at random.

dark states [22–24]. Internal dynamics can still move the system into states that do decay. If this decay is sufficiently rapid, however, it suppresses those transitions, a phenomenon known as the dissipative quantum Zeno effect [25–28]. Stronger loss therefore slows the escape of information from the protected states (Fig. 1c).

A long-lived state is not yet a usable memory, because each new input changes the same state that stores the past. A signal that initially depends on one input comes to depend on combinations of several inputs, such as their product. This mixing is useful for tasks that require relationships between inputs, but it reduces the capacity available for recalling single inputs [29–31]. Reading the memory has a cost as well, because resolving a weak quantum signal requires repeated preparation and measurement [10, 13, 32], and a long-lived memory is hard to read if its signal is faint. Previous work has established memory limits in classical reservoirs [2, 33, 34], explored input strength [35], and investigated how noise and protected states affect quantum information processing [36–39]. Closest to our work, Ref. [6] tunes the dissipation rate of the reservoir we study, Ref. [31] finds one trade-off between memory and nonlinearity across reservoir schemes, and Ref. [39] reads the output of a reservoir from a decoherence-free subspace to classify quantum states. The Zeno protection of dark states and the benefit of a weak input for linear memory are established. How protection, continued input, and measurement together determine the part of the past that a readout recovers is open.

Here we show that the form of the dissipation sets three properties of a memory (Fig. 1): how long a stored signal survives, whether and how strongly an input can be written into it, and how strong a signal is left to measure. First, the decay channel decides at equal total coupling whether stronger loss shortens or prolongs the memory lifetime. Independent decay shortens it, and collective decay prolongs it through dark states (Sec. 3.1). Second, mixing limits how strongly an input can be written. We prove that at a fixed input amplitude a stored signal combines on average $s\,t_{\mathrm{mem}}$ inputs, where the mixing fraction s is the part of the stored signal that one input disturbs and t_{mem} is the memory lifetime (Sec. 3.2). We also prove that an amplitude that falls as $1/\sqrt{t_{\mathrm{mem}}}$ removes this limit. The share of single inputs then converges to a value that depends on the input only through its squared amplitude times the lifetime, and this limit agrees with the five-spin results without a fitted parameter. The phases with which the spins couple to a collective channel decide whether an input reaches the dark

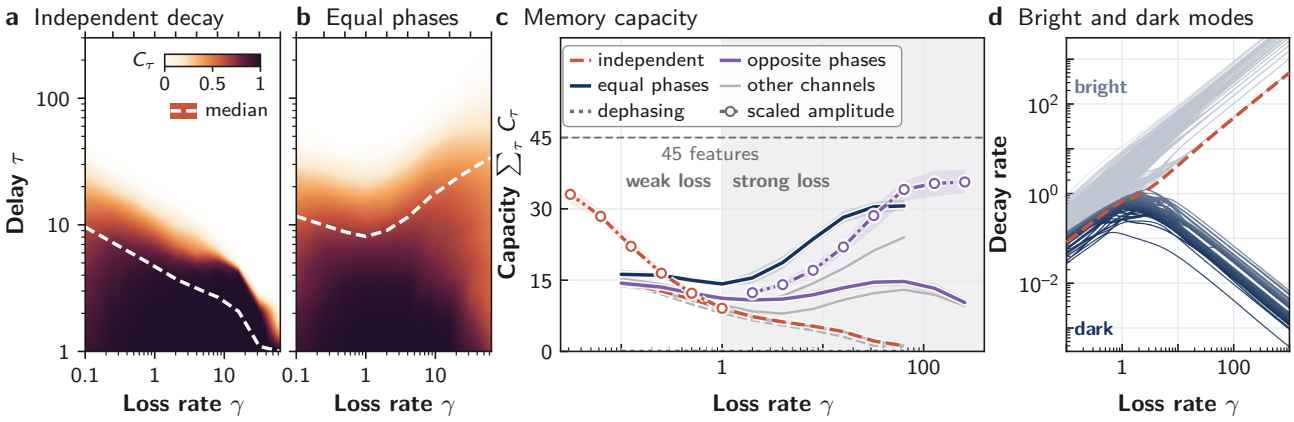


Figure 2: **The decay structure decides whether loss shortens or prolongs the memory. a,b** Recall of the input τ steps back: stronger loss pushes it to shorter delays with independent decay (a) and to longer delays with the equal-phase channel (b). **c** Linear memory capacity of all channels. Shading marks loss stronger than the couplings ($\gamma > 1$). **d** With the equal-phase channel the dark modes decay more slowly as the loss grows, while the slowest mode of independent decay (dashed) decays faster.

states at all (Sec. 3.3). Third, the state that the dissipation leaves behind sets the size of the signal that measurements must resolve. At strong loss a collective channel whose phases differ from those of the input field recovers the most memory at every measurement budget, and where the loss can be made weak, independent decay recovers more still (Sec. 3.4).

Together, these results connect the physical survival of information to what a readout recovers from it. The protection relies on quantum interference, whereas the two theorems hold for every reservoir that meets their conditions, classical ones included, and give the size of the trade-off between memory and nonlinearity. The form of the dissipation, the input amplitude and the measurement budget become joint design parameters, and the memory that a readout recovers becomes a quantity that can be engineered in physical systems that process information over time.

2 Model

The reservoir is the driven spin system of Ref. [6]. Five spins ($N = 5$) interact through random couplings J_{ij}, and in each step a random input u between 0 and 1 sets the strength of a transverse field for a time Δt:

$$H(u) = \sum_{i<j} J_{ij} X_i X_j + h \sum_i Z_i + h(1+u) \sum_i X_i, \tag{1}$$

with Pauli operators X_i, Z_i of spin i and a field strength h. To weaken the input at a fixed mean drive, we replace $h(1+u)$ by $h[\frac{3}{2} + \eta(u - \frac{1}{2})]$, where the amplitude η scales how far each input moves the field. The state ρ of the spins follows a Lindblad master equation [40, 41],

$$\dot{\rho} = -i[H(u), \rho] + \sum_k \Big(L_k \rho L_k^\dagger - \tfrac{1}{2}\{L_k^\dagger L_k, \rho\} \Big), \tag{2}$$

where the Hamiltonian drives the internal dynamics and each jump operator L_k, scaled by one loss rate γ, is one way the environment changes the state, such as removing an excitation or erasing a phase.

A linear readout of 45 features, the single-spin and two-spin expectation values, is trained to estimate the input from τ steps back. Its recall C_τ, between 0 and 1, is the squared correlation between the estimate and the input on held-out data, and the linear memory capacity $\sum_\tau C_\tau$ adds up the recall over the delays. One unit is one past input recalled perfectly, and the number of features bounds the capacity. The memory lifetime t_{mem} is the number of steps over which the variance of a stored signal survives (Appendix A.4). It is the only time scale of the main text, and Table 1 lists the others. Capacities in the text are sums over all delays with exact expectation values, as if every measurement were repeated infinitely often. The exceptions are Sec. 3.4, which counts shots, and Fig. 2c, which sums the first 300 delays of a simulated run. The capacities are numerical results for five spins with all-to-all couplings. The equal-coupling identity, the slowing of the dark modes, the mode-matching rule and the two theorems hold for any number of spins

and any Hamiltonian. How the signal scales with the number of spins is open, and reservoirs whose signals concentrate need exponentially many measurements [42].

We compare decay channels at equal total coupling $\mathrm{Tr}\, K = \sum_k \mathrm{Tr}\, L_k^\dagger L_k$ and hold the Hamiltonian fixed to isolate the effect of the form of the dissipation. The generator of Eq. (2) has 4^N modes, patterns of change in the state that each decay at their own rate. Equal total coupling fixes their mean rate for every Hamiltonian, but not the slow modes, which hold a stored signal longest and become slower or faster with the channel (Appendix A.2). Table 2 lists the eight channels.

We discuss four of them (Fig. 1b). With independent decay every spin loses its excitation on its own. With collective decay all spins lose it through one shared channel, with equal phases or with opposite phases on spins 4 and 5, and we refer to the equal-phase and the opposite-phase channel. With dephasing every spin loses its phase. The channels differ in the kernel of the dissipator, the operators that the dissipation alone leaves unchanged, although the Hamiltonian can still turn them or couple them to decaying ones. For independent decay the kernel is the ground state alone. A collective channel cannot empty its dark states, so its kernel contains every operator on the dark subspace. There are $\binom{N}{\lfloor N/2 \rfloor}$ dark states, ten for five spins, and their number grows exponentially with the number of spins. Dephasing keeps the populations, the probabilities of the spin configurations, and erases the coherences between them. Every step then leaves the fully mixed state unchanged whatever the input, the driven reservoir relaxes to this state, and dephasing stores nothing.

3 Results

3.1 The decay channel sets the memory lifetime

Increasing the loss separates the channels into two groups (Fig. 2). With independent decay, recall retreats to shorter delays and the linear memory capacity falls steadily (Fig. 2a,c). With the equal-phase channel the capacity first dips and then doubles, and at strong loss inputs fifty steps back are still recalled (Fig. 2b,c). The other channels follow one of the two (Fig. 2c, grey). Exchange and gain/loss follow independent decay, two collective channels follow the equal-phase channel, and pairwise decay follows the opposite-phase channel. The decay rates explain the split, because without dark states the slowest rate grows in proportion to γ and the memory lifetime shrinks, while with dark states the slow rates fall as $1/\gamma$ and the lifetime grows (Fig. 2d).

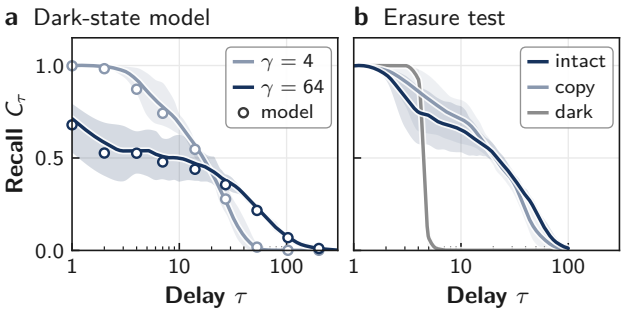


Figure 3: **Dark states hold the memory.** Equal-phase channel. **a** A model of the dark states alone (circles) reproduces the recall of the full reservoir (lines). **b** Erasing what the dark states know about the inputs destroys the memory (dark), while erasing a rotated copy of them leaves most of it (copy).

Stronger collective decay prolongs the memory through the quantum Zeno effect [25, 43]. In the two-spin picture of Fig. 1c, the dark state $|-\rangle \propto |10\rangle - |01\rangle$ can decay only through the bright state $|+\rangle \propto |10\rangle + |01\rangle$, and strong loss suppresses this escape. Its rate falls as Ω^2/γ, where Ω couples $|+\rangle$ and $|-\rangle$, as the slow mode of an overdamped oscillator settles more slowly the stronger the friction.

For any number of spins the protection follows from the strong-loss expansion at fixed system size. At strong loss the reservoir follows its Hamiltonian restricted to the dark states, which turns the stored state without making it decay. Decay enters at the next order, when the Hamiltonian briefly lifts a dark state into a bright one [26, 28]. One step is then a rotation plus a loss that falls as $1/\gamma$ (Eq. (A2)), a nearly orthogonal map with a small decay rate, which is the classic recipe for a long memory in a network [33]. Calculations for three to six spins show the same decay out of the dark states with different prefactors, and with the equal-phase channel the capacity also grows with the loss at four and six spins (Appendix E.1).

The dark states hold the memory, as two tests show (Appendix E.1). A model of the ten dark states alone, with no fitted parameter, reproduces the recall of the full reservoir at strong

loss (Fig. 3a). Replacing the dark-state components of the state by their time averages after every step, which removes their dependence on the inputs, destroys all memory beyond a few steps, while the same operation on a randomly rotated set of equal size leaves most of it (Fig. 3b). Residual local loss limits the lifetime, and for the equal-phase channel recalling inputs over T steps requires a ratio of collective to local loss rates of at least $4.2\,T^2$ (Appendix E.6), so that a ratio of 10^4 protects 49 steps. For superconducting qubits coupled to a waveguide, a four-qubit dark state outlives the decay of a single qubit into the waveguide by more than two orders of magnitude [44], and qubit arrays reach a cooperativity of 172 [45].

Dark states protect a stored signal against decay but not against the inputs that follow it. With the opposite-phase channel the dark states are just as protected as with the equal-phase channel, yet the capacity returns only to its weak-loss level (Fig. 2c).

3.2 New inputs mix the stored signal

Each new input turns the stored state by a small angle that depends on the input (Fig. 1d), so every newer input acts as noise on an older one. Over the memory lifetime the stored direction takes a random walk of small turns $\pm\theta$ and spreads by about $\theta\sqrt{t_{\mathrm{mem}}}$. Once the spread reaches about a radian, an old input survives only multiplied by the inputs that followed it, and a linear readout cannot turn such products back into single inputs.

Each part of a stored signal depends on one past input or on several jointly, and we denote the number of distinct inputs that it depends on by D. A function of u_{t-3} alone has $D = 1$, even if it is nonlinear in u_{t-3}, and $u_{t-3}u_{t-5}$ has $D = 2$. The information processing capacity [29, 46] also counts how well the readout reconstructs nonlinear functions of past inputs, and sorting its contributions by D separates single inputs from combinations of inputs (Appendix A.4). The mixing fraction s, between 0 and 1, is the part of a stored signal that one input turns differently from the average input, and for small turns it grows as the square of the turn. Our first theorem concerns the covariance of the state across input histories and assumes that one direction of this covariance outlives all others, even when many slow components of the state share it. For reservoirs that leak slowly and are mixed by their inputs, it states that a part of the stored signal involves on average $s\,t_{\mathrm{mem}}$ inputs.

> **Theorem 1** (Mixing law, made precise in Appendix B)**.** Let a reservoir be reset with probability p per step, so that $t_{\mathrm{mem}} \sim 1/(2p)$ as $p \to 0$, and otherwise turned by an input-dependent rotation. Let the average rotation shrink every state, let the rotations leave no stored covariance unchanged on average except the isotropic one, and let one input change the stationary state (conditions (1) to (3) of Appendix B). Pick a part of the stored signal with probability equal to its share of the capacity, and let D be the number of distinct inputs it depends on. As $p \to 0$, D/t_{mem} becomes exponentially distributed with mean s, and
>
> $$\mathbb{E}D = s\,t_{\mathrm{mem}} + O(1).$$

The law describes a race in which, in each step, a fraction s of the stored signal picks up one more input while a fraction $1/t_{\mathrm{mem}}$ of it is lost. The number of inputs then follows a geometric distribution, which becomes exponential for long memories.

The law is proved for the reset model in any dimension, and nothing in the proof is quantum. Classical state-affine reservoirs, whose state is updated by an input-dependent linear map plus an offset [47], obey it under the same conditions. The proof carries over to physical channels without a reset, with the slow decay in place of the reset, when a single direction of the stored covariance is slow, whereas several slow directions keep distinct mixing scales (Appendix B.3). For three spins, interval arithmetic certifies the conditions (Appendix D), and for five spins the recursion of Eq. (A3) gives the distribution of D exactly, which we compare with the law (Fig. 4c).

At a fixed amplitude, $\eta = 1$, one input disturbs the stored state only slightly, $s \approx 0.02$. At short memories the mean of D is slightly below two, because the fast response of the reservoir combines the latest inputs within a few steps. As the memory grows, the mean of D follows $s\,t_{\mathrm{mem}}$ to within about two inputs, up to 23 inputs at the longest memory, and the distribution of D approaches the exponential law (Appendix B.3). Dissipation enters the law through its two factors, the memory lifetime and the mixing fraction, and

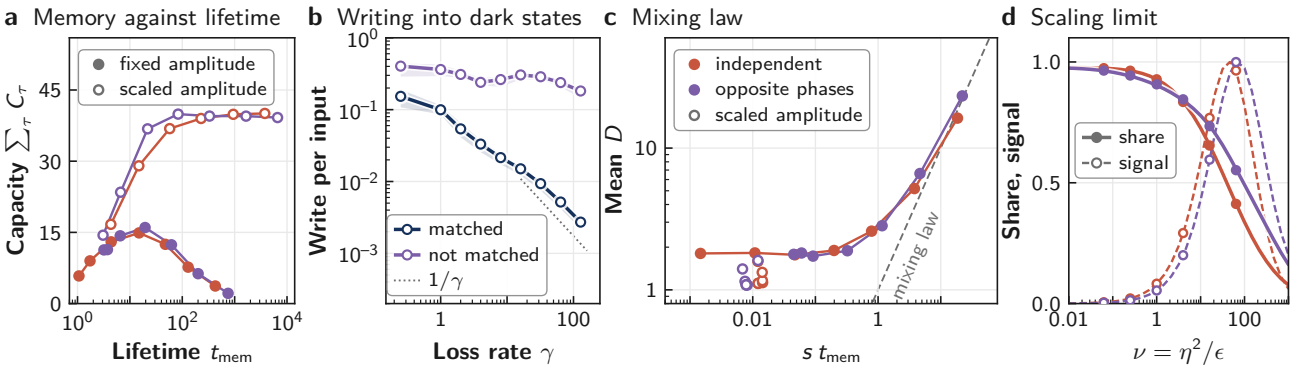


Figure 4: **New inputs mix the stored signal, and a scaled amplitude limits the mixing.** **a** At a fixed input amplitude the capacity, summed over all delays, peaks and then falls as the memory grows, and a scaled amplitude holds it at 40 units. **b** An input that is mode-matched to the equal-phase channel writes less into the dark states as the loss grows, while an input that is not mode-matched keeps writing. **c** The number of inputs mixed into a stored signal follows the mixing law, and a scaled amplitude keeps it near one. **d** The scaling limit (lines) against the five-spin recursion (markers), without a fitted parameter. Solid lines and filled markers show the share of single inputs in the capacity, dashed lines and open markers the signal they leave in the features, scaled to its maximum.

their product is the size of the trade-off between memory and nonlinearity [31].

At a fixed amplitude, lengthening the memory beyond $s\,t_{\mathrm{mem}} \approx 0.3$ lowers the linear memory capacity (Fig. 4a). With the opposite-phase channel it peaks at 16 units and falls to 2 units while the lifetime grows from 20 to 750 steps, and independent decay behaves in the same way (Appendix E.3). The stored signal has moved into products of inputs, and the share of the capacity carried by single inputs falls from 0.37 to 0.05. Single inputs remain readable only if the mixing fraction shrinks as the memory lifetime grows, which the phases of the decay channel and the amplitude of the input both achieve.

3.3 Channel phases and input amplitude control the mixing

The phases of the collective channel decide whether an input reaches the dark states. Spin i couples to the shared channel with the sign $c_i = \pm 1$, and the dark states are annihilated by the collective lowering operator $\sum_i c_i \sigma_i^-$. An input field is mode-matched to the channel if its strength on each spin is proportional to c_i, so that it drives the collective mode through which the spins decay. A mode-matched field $\sum_i c_i X_i$ is the lowering operator plus its adjoint, so it cannot couple two dark states, for any number of spins and any Hamiltonian. It writes into them only through brief excursions into bright states, with a strength that falls as $1/\gamma$, while the part of any other field that is not mode-matched writes into them directly (Fig. 4b, Appendix E.2). The uniform input of Eq. (1) is mode-matched to the equal-phase channel. Its mixing fraction shrinks as the loss grows, and this channel reaches twice the capacity of the other two at a fixed amplitude (Table 3). With the opposite-phase channel the same input is not mode-matched, writes directly into the dark states and mixes them at every loss rate. This is the generic case, because a phase error opens a direct coupling within the dark subspace (Appendix E.2).

The input amplitude η controls the mixing for every channel. A weak input is known to favour linear memory [30, 31, 35], because it removes the products that the fast response forms within a few steps. The mixing law adds what a long memory requires. The mixing fraction grows as the square of the amplitude, so $s\,t_{\mathrm{mem}}$ stays fixed only if the amplitude falls as $1/\sqrt{t_{\mathrm{mem}}}$, which we refer to as a scaled amplitude. Theorem 1 shows that the scaling is necessary, since at a fixed amplitude the share of single inputs vanishes as the memory grows. Our second theorem shows that it is sufficient.

Theorem 2 (Scaling limit, made precise in Appendix B.4)**.** Let one step of a reservoir be a rotation, a leak of small rate ϵ and an input-dependent turn of amplitude η, under the decay condition of Appendix B.4. As $\epsilon \to 0$ with $\nu = \eta^2/\epsilon$ fixed, the stored signal and the

distribution of D converge to the solution of a linear problem that depends on the input only through ν. The mean of D stays finite, and single inputs keep a fixed share of the capacity. As $\nu \to \infty$ the limit joins the mixing law, with $s\,t_{\mathrm{mem}}$ proportional to ν.

A scaled amplitude holds ν fixed, so the memory lifetime grows without limit while the composition of the stored signal stays the same. The leak is $\epsilon = \gamma$ for independent decay and $\epsilon = 1/\gamma$ for the dark states, and the limit follows from the Hamiltonian and the jump operators without a fitted parameter. It reproduces the five-spin recursion over three decades of ν, with the share of single inputs agreeing to within 0.01 (Fig. 4d, Appendix E.3). At $\nu = 1$ single inputs carry about 0.9 of the capacity and the linear memory capacity is 40 units for both channels, the level at which the recursion settles as the lifetime grows (Fig. 4a). One qubit and three spins are special cases, with a closed form and with certified bounds (Appendices C and D). The three properties of the memory appear in the limit as three objects: the leak, the first-order response to the input field, which vanishes on the dark states for a mode-matched input, and the change that an input makes to the stationary state (Appendix B.4).

In the five-spin reservoir $\nu = 1$ is an amplitude of $1/\sqrt{\gamma}$ for the opposite-phase channel and of $\sqrt{\gamma}$ for independent decay (Fig. 2c, circles). Compared with the best loss rate at a fixed amplitude, the linear memory capacity rises from about 15 to 40 units. Both effects of a weak input lower the mean of D, and the mixing over the lifetime accounts for almost all of the reduction once the lifetime exceeds a few hundred steps (Appendix E.3). The capacities of all values of D add up to the number of features [29], so the gain moves capacity from combinations of inputs to single inputs. A scaled amplitude is the choice for recall, and a task that needs products of inputs calls for a larger amplitude, with $s\,t_{\mathrm{mem}}$ giving the number of inputs that a stored signal then combines. A long memory needs no dark states, since weak independent decay with a scaled amplitude reaches one too. Dark states add a memory that stays long under strong dissipation and grows with the loss.

The scaled amplitude is an upper bound on the input and not an optimum. With an exact readout the capacity keeps rising as the input weakens, towards the number of features (Appendix E.4), because a weaker input mixes less and a fainter signal costs nothing without measurement noise. The lifetime bounds the amplitude only from above, $s\,t_{\mathrm{mem}} \lesssim 1$. Theorem 2 quantifies what a weak input gives up (Fig. 4d). The signal that single inputs leave in the features vanishes in proportion to ν for weak inputs and is largest where a stored signal combines a few inputs and single inputs carry about half of the capacity (Appendix E.3). The limit depends on the amplitude only through ν, so the amplitude of the largest signal also falls as $1/\sqrt{t_{\mathrm{mem}}}$.

For three spins with the opposite-phase channel, computer-checked interval bounds confirm both regimes at strong loss. At a fixed amplitude single inputs lose their share of the capacity, and with an amplitude $\eta_0/\sqrt{\gamma}$ they keep it (Appendix D). The stored signal resides in coherences between the energy levels of the Hamiltonian restricted to the dark states. Coherence shapes the performance of spin reservoirs [48], and erasing these coherences after every input reduces the stored signal as $1/\gamma^2$. A weaker input leaves a fainter signal, which is free with an exact readout and costly with measurements.

3.4 The stationary state sets the measurable signal

Exact expectation values remove the cost of resolving weak signals. A memory with a scaled amplitude grows longer without growing larger, because the capacity cannot exceed the number of features and each delay is recalled more faintly. With projective readout each feature is estimated from n shots per measurement setting and step, and its statistical error must stay below the signal, the variation of the feature with the inputs (Appendix A.5). The capacity recovered from a finite number of shots is the counterpart, for the recall of past inputs, of the resolvable expressive capacity of Ref. [32].

The state that the dissipation leaves behind sets the size of the signal. An input acts on the stationary state ρ_* through the commutator with its Hamiltonian, and only the part of this change that lies on the slow modes is remembered. Its size is the linear response of the stationary state to the input field, which we call the stationary response, in the spirit of the link between memory and

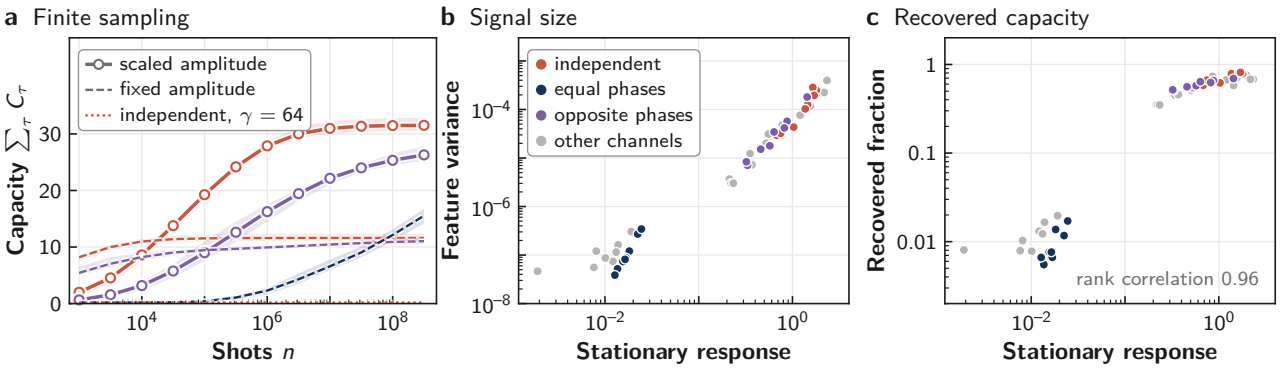


Figure 5: **Finite sampling limits the memory that can be read out. a** Capacity recovered against the number of shots per setting and step, with replayed inputs, for independent decay at weak loss ($\gamma = 1/32$) and the two collective channels at strong loss ($\gamma = 64$). **b** The signal grows with the stationary response. **c** So does the fraction of the exact capacity that 10^6 shots recover. Points in b and c are single reservoirs at loss rates that give the same lifetime, with the full input history.

optical absorption of Ref. [9]. When all channels are compared at loss rates that give the same lifetime, the feature variance of the seven channels that respond follows the stationary response over four orders of magnitude (Fig. 5b). The fraction of the exact capacity that 10^6 shots recover follows it as well, with a rank correlation of 0.96 over all reservoirs (Fig. 5c). Collective decay with equal phases leaves a polarized state that the uniform input moves only through bright states. It has the largest exact capacity of all channels and recovers 1% of it. Reversing the coupling phase of two spins lets the input move the state directly (Sec. 3.3), and the opposite-phase channel recovers 60% of its exact capacity, against 73% for independent decay (Appendix E.5). The exact capacity does not predict what measurements recover, and the stationary response predicts which fraction of it they do.

The form of the dissipation matters most where the loss is strong (Fig. 5a). At $\gamma = 64$ and a fixed amplitude, 10^4 shots recover 8 units with the opposite-phase channel and only the floor of no recall with the equal-phase channel or with independent decay (Table 3). A collective channel that is not mode-matched to the input is the right choice when strong loss is given, as for emitters that radiate into a common mode [22–24]. When the loss can be made weak, independent decay recovers more, 11 units from the same number of shots.

A short memory needs few shots, and a scaled amplitude helps only at larger numbers. Its capacity overtakes that of the fixed amplitude between 10^4 and 10^5 shots (Appendix E.5).

Every shot replays the recent inputs before the measurement [13], and the replayed window grows with the memory lifetime. When every replayed step is counted (Appendix A.5), a budget of 10^9 reservoir steps per output step gives the three channels the same capacity at a fixed amplitude, each at its best loss rate (Table 3). A collective channel that is not mode-matched holds at a hundredfold stronger loss what independent decay holds at weak loss. With a scaled amplitude, weak independent decay recovers the most, 34 units against 21 with the opposite-phase channel.

With a finite number of shots a best amplitude exists. It falls as a power of the lifetime with a fitted exponent close to the $-1/2$ of the scaled amplitude (Appendix E.4). The lifetime bounds the input amplitude from above, and the measurement budget bounds it from below.

4 Conclusion

A useful quantum reservoir must keep a readable signal of the past while it accepts new inputs. We have shown that the form of the dissipation sets three properties of this memory. Its lifetime depends on the decay channel, because independent decay shortens the memory as the loss grows, whereas strong collective decay prolongs it by holding the stored signal in dark states through the quantum Zeno effect.

Writing is limited by mixing, because a longer memory leaves more time for later inputs to mix with the stored signal, and the mixing law gives the number of inputs that mix. The input ampli-

tude must fall as the inverse square root of the lifetime, and the scaling limit shows that this is enough. The share of single inputs then follows from a linear problem that reproduces the simulations without a fitted parameter, and in the spin reservoirs studied the linear memory capacity rises from about 15 to 40 of 45 units with an exact readout. The phases of a collective channel decide whether an input reaches the dark states at all.

The signal must be strong enough to resolve, and here the equal-phase channel holds the largest memory and leaves the faintest signal. At strong loss a collective channel whose phases differ from those of the input field keeps a readable memory where independent decay keeps none. Where the loss can be made weak, independent decay recovers the most because its signals are stronger.

The environment, the input and the readout have to be designed together, and three rules summarize the design:

(1) *Lifetime.* Collective decay protects a memory in dark states. Stronger loss then makes it last longer, as long as local loss stays weak.

(2) *Writing.* The input must reach the dark states, and an input with the phases of the decay channel reaches them ever more weakly as the loss grows. The longer the memory lasts, the weaker each input must be, or new inputs mix with the old ones.

(3) *Signal.* What can be read out depends on how strongly the inputs move the state that the dissipation leaves behind. At strong loss a collective channel whose phases differ from those of the input recovers the most.

Acknowledgments

The authors thank the Quantum Applications and Research Laboratory and the Mobile and Distributed Systems Group at Ludwig-Maximilians-Universität München for institutional support.

Table 1: Quantities of the text. Times are in input steps of length Δt.

t_{mem}	memory lifetime: how long the variance of a stored signal survives, whatever inputs it combines (Appendix A.4)
t_{f}	recall time: how long one input stays linearly readable on its own (Appendix A.3)
t_{rel}	relaxation time: the recall time with the input held at its mean value (Appendix A.5)
τ_{90}	replayed range: the first τ_{90} delays hold 90% of the exact recall
D	number of distinct past inputs that a part of the stored signal depends on
s	mixing fraction: the part of a stored signal that one input turns differently from the average input
ν	squared input amplitude per leak, η^2/ϵ, which a scaled amplitude holds fixed (Appendix B.4)

Data availability

The code, the data behind every figure, and the derivation and certificates of the three-spin limits are at `https://github.com/eybmits/what-a-dissipative-quantum-system-remembers`.

Author contributions

M.B. contributed to conception, methodology, software, analysis, visualization and writing, I.F. and J.W. to methodology, interpretation and editing, and C.L.-P. and J.S. to supervision and conceptual development. OpenAI Codex and Claude Code (Anthropic) assisted with manuscript preparation, checks of derivations and references, and simulation, verification and plotting code. The authors reviewed the outputs and take responsibility for the final content.

A Methods

This appendix specifies the five-spin simulations and defines the quantities of the main text, which Table 1 lists.

A.1 Reservoir, readout and capacity

The Hamiltonian of Eq. (1) uses $h = 1/2$ and couplings J_{ij} drawn uniformly from $[-1, 1]$, so the transverse field runs from h to $2h$. Each input u is drawn uniformly from $[0, 1]$ and held for a step of

length $\Delta t = 1/2$. We integrate Eq. (2) in the 1024-dimensional basis of Pauli operators and evaluate the step map $E(u) = e^{\Delta t \mathcal{L}(u)}$ by Chebyshev interpolation in u at ten nodes. Each reservoir draws its couplings and its input stream from one seed, and every channel uses the same eight seeds. After a washout of $\max(800, 60\gamma)$ steps, at least four times the recall time, we record the 45 features, the expectation values of X_i, Y_i, Z_i and of X_iX_j, Y_iY_j, Z_iZ_j, for 2000 training, 400 validation and 1600 test inputs. The readout is a least-squares fit with an intercept on the training and validation inputs, and C_τ is the squared correlation of its estimate with $u_{t-\tau}$ on the test inputs. The linear memory capacity sums C_τ over delays 1 to 300. A delay without any recall reads about $1/1600$ on the test inputs, so the sum has a floor of about 0.2 units. Figures show the mean of the eight reservoirs, with bands from the 16th to the 84th percentile.

A.2 Decay channels and equal coupling

Table 2 lists the jump operators of the eight decay channels. Their weights give every channel the same total coupling $\operatorname{Tr} K = N 2^{N-1}\gamma = 80\gamma$, which fixes the mean decay rate. For a traceless jump L the superoperator trace of its dissipator is $-2^N \operatorname{Tr} L^\dagger L$, while the commutator with any Hamiltonian is traceless, so the 4^N eigenvalues λ of the generator obey

$$\frac{1}{4^N} \sum \operatorname{Re} \lambda = -\frac{\operatorname{Tr} K}{2^N} \tag{A1}$$

for every channel and every Hamiltonian. The kernel of the dissipator is larger than the set of dark states: the ten dark states of a collective channel carry $10 \times 10 = 100$ operators.

A.3 Slow modes, recall time and the strong-loss generator

The input-averaged step $\bar{E} = \mathbb{E}_u E(u)$ is computed by six-point Gauss–Legendre quadrature. Its eigenvalues λ_j give decay rates $\kappa_j = -\ln|\lambda_j|/\Delta t$. If the kernel of the dissipator has dimension n, the slow modes are the $n-1$ slowest decaying ones besides the stationary state, and the recall time is $t_\mathrm{f} = 1/(\Delta t\, \bar{\kappa})$, with $\bar{\kappa}$ their mean rate. For $n = 1$ we use the slowest mode. The averaged step carries the part of a stored signal that still depends on a single input. Its mean rate adds decay and mixing to second order in the amplitude,

Table 2: Decay channels at five spins, each with $\operatorname{Tr} K = 80\gamma$, and the number of operators that the dissipator leaves unchanged. The operators $\sigma_i^{\mp}$ lower and raise spin i, with $i < j$ for pairwise decay, $i \neq j$ for exchange and $c = (1, 1, 1, -1, -1)$.

channel	jump operators	kernel dimension
independent	$\sqrt{\gamma}\,\sigma_i^-$	1
equal phases	$\sqrt{\gamma}\sum_i \sigma_i^-$	100
opposite phases	$\sqrt{\gamma}\sum_i c_i\sigma_i^-$	100
two channels	$\sqrt{\gamma}\sum_{i\le 3}\sigma_i^-$, $\sqrt{\gamma}\sum_{i>3}\sigma_i^-$	36
pairwise	$\sqrt{\gamma}\,\sigma_i^-\sigma_j^-$	36
exchange	$\sqrt{0.4\gamma}\,\sigma_i^-$, $\sqrt{0.3\gamma}\,\sigma_i^-\sigma_j^+$	1
gain/loss	$\sqrt{\gamma/1.3}\,\sigma_i^-$, $\sqrt{0.3\gamma/1.3}\,\sigma_i^+$	1
dephasing	$\sqrt{\gamma/2}\,Z_i$	32

$\bar{\kappa} = \bar{\kappa}_0 + c\,\eta^2 + O(\eta^4)$, because the inputs are distributed symmetrically about their mean and the term of $\bar{E}$ linear in η vanishes. Appendix E.2 tests this sum numerically. Figure 2d shows instead the decay rates $-\operatorname{Re}\lambda$ of the generator at the mean input, $u = 1/2$. With a collective channel we call the $n-1$ slowest of them dark and all faster ones bright.

Let $\mathcal{D}$ be the unit-rate dissipator, $\mathcal{P} = \lim_{t\to\infty} e^{t\mathcal{D}}$ the projector onto its kernel and $\mathcal{S} = (\mathcal{D}+\mathcal{P})^{-1} - \mathcal{P}$ its inverse on the decaying operators. At strong loss, eliminating the decaying operators to second order [26, 28] gives the generator on the kernel,

$$\mathcal{G}(u) = \mathcal{P}\mathcal{H}(u)\mathcal{P} - \frac{1}{\gamma}\,\mathcal{P}\mathcal{H}(u)\mathcal{S}\mathcal{H}(u)\mathcal{P}, \tag{A2}$$

with $\mathcal{H} = -i[H, \cdot]$. The first term is the Hamiltonian restricted to the kernel, which turns its operators, and the second is the loss through brief excursions into bright states.

The scaled amplitude is $\eta = \min(1, \gamma^{-1/2})$ for the opposite-phase channel and $\eta = \min(1, \gamma^{1/2})$ for independent decay, which is $\nu = 1$ in Appendix B.4.

A.4 Number of inputs mixed

The order D of a stored signal is the number of distinct past inputs it depends on (Sec. 3.2). Precisely, for independent inputs, every function of the past inputs with finite variance splits uniquely into orthogonal parts, each of which depends on a fixed finite set of inputs. The order of a part is the size of its set, not a polynomial degree. For signs $u_t = \pm 1$ the parts are the products $\prod_{j\in J} u_{t-j}$, and

for continuous inputs one uses orthogonal polynomials in each input. We collect the traceless coefficients of the state in an orthonormal operator basis in a vector r, on which quantum channels act affinely [49, 50], $r_t = A(u_t)r_{t-1} + b(u_t)$. All means and covariances are taken over input histories. Let $\bar{A} = \mathbb{E}A$ be the mean step, m the stationary mean and $\xi = A(u)m + b(u) - m$ the innovation, the change that one input writes, with covariance V. With $\mathcal{B}X = \bar{A}X\bar{A}^\mathsf{T}$, $\mathcal{T}X = \mathbb{E}[AXA^\mathsf{T}]$ and $\mathcal{K} = \mathcal{T} - \mathcal{B}$, the covariance G_k carried by products of order k obeys

$$G_1 = (1-\mathcal{B})^{-1}V, \qquad G_{k+1} = (1-\mathcal{B})^{-1}\mathcal{K}G_k. \tag{A3}$$

The innovation adds the current input alone. The mean step carries a stored product along without adding an input. The fluctuating part $\mathcal{K}$ of the step multiplies it by the current input, which raises its order by one. The G_k add up to the stationary covariance Σ. A readout of features $y = Mr$ has the capacity $\mathrm{Tr}[(M\Sigma M^\mathsf{T})^+ M G_k M^\mathsf{T}]$ of order k, and these capacities add up to the rank of the readout [29]. Divided by their sum, they give the share of each order, the distribution of D in Theorem 1. For one observable with coordinates a the share of order k is $a^\mathsf{T}G_k a/a^\mathsf{T}\Sigma a$, and the share of single past inputs is $a^\mathsf{T}\bar{A}G_1\bar{A}^\mathsf{T}a/a^\mathsf{T}\Sigma a$.

The memory lifetime and the mixing fraction come from the slowest mode of the covariance step $\mathcal{T}$: its eigenvalue is $1 - 1/t_{\mathrm{mem}}$, and with its right eigenvector Q and left eigenvector ℓ the mixing fraction is $s = \ell(\mathcal{K}Q)/\ell(Q)$, as in Appendix B.3. For five spins we evaluate Eq. (A3) on the 1023 traceless Pauli coordinates (Appendix E.3).

A.5 Finite sampling and the measurement budget

A quantum measurement gives a random outcome, so every feature is an average over n shots per measurement setting and step. We use three settings, with all spins read along X, Y or Z. For each step we form the outcome distributions of the three settings and draw n shots from each. The readout is fitted as above, with a ridge chosen from fourteen values between 0 and 10^4 on the validation inputs. In Sec. 3.4 the distributions come from the replay: each output step replays its last τ_{90} inputs from $|00000\rangle$, where the first τ_{90} delays hold 90% of the exact recall. We count the budget B in reservoir steps, shots times replayed steps summed over the three settings, per output step, so that $n = B/(3\tau_{90})$, and we also run $n = 10^3$ to $10^{8.5}$ in half decades. Only Fig. 5b,c and Table 4 use the state with the full history. Their loss rates are matched to a relaxation time $t_{\mathrm{rel}} = 1/(\Delta t\,\bar{\kappa}_0)$ of 32 input steps, where $\bar{\kappa}_0$ is the mean decay rate of the slow modes of the generator at the mean input, without the mixing by the inputs (Appendix E.5).

B The mixing law and the scaling limit

This appendix proves the two theorems and extends the mixing law to physical channels without a reset.

B.1 Setting and statement

We describe the reservoir by the n traceless coordinates r of its state in an orthonormal basis of operators, the Bloch vector for a qubit. A unitary acts on r as a rotation. In each step the reservoir is reset to a fixed state with a small probability p and is otherwise rotated by $O(u)$:

$$r_t = (1-p)\,O(u_t)\,r_{t-1} + p\,b. \tag{B1}$$

Here b represents the reset state. Equation (B1) averages over whether the reset occurred, so the only remaining randomness is the input history, with independent and identically distributed inputs u_t. A stored signal loses a fraction $1-(1-p)^2 = 2p-p^2$ of its variance per step, so its memory lifetime, defined by the slow eigenvalue $1 - 1/t_{\mathrm{mem}}$ of the covariance step (Appendix A.4), is $t_{\mathrm{mem}} = 1/(2p-p^2)$, or $1/(2p)$ to leading order. Three conditions make the law hold:

(1) *The average step forgets.* The average rotation $A = \mathbb{E}O$ shrinks every vector: its spectral radius is below one.

(2) *The inputs mix.* The only covariance matrices that the random rotations leave unchanged on average, $\mathbb{E}[OXO^\mathsf{T}] = X$, are multiples of the identity.

(3) *The input writes.* One input changes the state: $\mathbb{E}\|(O-A)\mu\|^2 > 0$ with $\mu = (1-A)^{-1}b$.

The first condition means that only products of inputs can be stored for long, the second that a stored signal spreads evenly over all directions,

and the third that there is something to store. The mixing fraction

$$s = 1 - \frac{\|A\|_F^2}{n} \tag{B2}$$

says how strongly one input turns the stored state. For a unit vector x of random orientation, $\mathbb{E}\|(O-A)x\|^2 = s$: one input turns a fraction s of a stored signal away from the average rotation. The fraction lies between 0 and 1, the first condition makes it positive, and for small turns θ it grows as θ^2.

Mixing law, precise form. Assume conditions (1)–(3) and take any fixed linear readout. Pick a part of the stored signal with probability equal to its share of the readout's capacity, and let D be its order, the number of distinct inputs it depends on. As $p \to 0$, the ratio D/t_{mem} becomes exponentially distributed with mean s, and

$$\mathbb{E}D = s\, t_{\text{mem}} + O(1).$$

Equivalently, products of more than $x\, t_{\text{mem}}$ inputs carry the share $e^{-x/s}$ of the capacity.

B.2 Proof

The proof makes the race of Sec. 3.2 exact: on average $1/s$ steps pass before the next input joins a product, which meanwhile loses a fraction $1/(s\, t_{\text{mem}})$ of its weight. The proof works with covariance matrices. Appendix A.4 splits the stationary covariance Σ of the state into the parts G_k carried by products of k inputs and gives their recursion, Eq. (A3): the mean step $\mathcal{B}$ carries a stored product along, and the fluctuating part $\mathcal{K}$ of the step adds one input to it. In the reset model the mean step is $\mathcal{B}X = (1-p)^2 AXA^\mathsf{T}$, the full step $\mathcal{T}X = (1-p)^2\,\mathbb{E}[OXO^\mathsf{T}]$, and $\mathcal{K} = \mathcal{T} - \mathcal{B}$. We write $\mathcal{B}_0$, $\mathcal{T}_0$ and $\mathcal{K}_0$ for the same maps at $p = 0$.

Step 1: the stored signal is small and spread evenly. The stationary mean is $m = p\,[1-(1-p)A]^{-1}b = p\mu + O(p^2)$, so one input writes the change $(1-p)(O-A)m$, of size p. Its covariance V has the trace $p^2 v + O(p^3)$ with $v = \mathbb{E}\|(O-A)\mu\|^2$, which is positive by condition (3). The random rotations preserve the trace of a covariance and, by condition (2), leave only the identity unchanged. Since they cannot enlarge a covariance, the identity is a simple fixed point with the projector $X \mapsto (\operatorname{Tr} X/n)\,1$. On every other direction $(1-\mathcal{T})^{-1}$ stays bounded as $p \to 0$, while on the identity it equals $1/[1-(1-p)^2] \approx 1/(2p)$. The written covariance therefore piles up along the identity:

$$\Sigma = (1-\mathcal{T})^{-1}V = \frac{pv}{2n}\,1 + O(p^2). \tag{B3}$$

Step 2: few inputs carry a vanishing share. The mean step forgets by condition (1), so $(1-\mathcal{B})^{-1}$ stays bounded as $p \to 0$. The recursion then makes $G_1 = (1-\mathcal{B})^{-1}V$ of order p^2, and the same holds for every G_k with fixed k. Measured against Σ, which is of order p and has full rank, any fixed set of orders carries a share of order p of the capacity of any readout. A long memory consists of products of many inputs.

Step 3: counting the inputs. Call $\mathcal{F} = (1-\mathcal{B})^{-1}\mathcal{K}$ the order transfer, so that $G_{k+1} = \mathcal{F}G_k$. At $p = 0$ it leaves the identity unchanged, because $\mathcal{K}_0 1 = 1 - AA^\mathsf{T} = (1-\mathcal{B}_0)1$. Its matching left eigenvector is the weight

$$\ell(X) = \operatorname{Tr}[(1 - A^\mathsf{T}A)X], \qquad \ell(1) = ns, \tag{B4}$$

since $\ell(\mathcal{F}_0 X) = \operatorname{Tr}\mathcal{K}_0 X = \ell(X)$ when the rotations preserve the trace. For small p the eigenvalue one moves to

$$1 - \frac{2}{s}\,p + O(p^2)\,: \tag{B5}$$

each step up in order costs a fraction $2p/s = 1/(s\, t_{\text{mem}})$ of the weight, as in the race. To see this, differentiate $(1-\mathcal{B})\mathcal{F} = \mathcal{K}$ at $p = 0$ and apply it to the identity, which gives $(1-\mathcal{B}_0)\mathcal{F}_0' 1 = -2\cdot 1$. Its trace is the weight $\ell(\mathcal{F}_0' 1) = -2n$, and dividing by $\ell(1) = ns$ gives the rate $2/s$.

To turn the eigenvalue into a distribution, weigh order k with $e^{-\lambda p k}$. Summing the recursion gives $\sum_k z^k G_k = z(1-z\mathcal{F})^{-1}G_1$, and as $z = e^{-\lambda p}$ approaches one, only the eigenvalue of Eq. (B5) makes the inverse large. With $\ell(G_1) = p^2 v + O(p^3)$ and $\beta = 2/s$,

$$\sum_{k\geq 1} e^{-\lambda p k} G_k = \frac{pv}{2n}\,\frac{\beta}{\lambda+\beta}\,1 + O(p^2). \tag{B6}$$

Divided by Eq. (B3), this is $\beta/(\lambda+\beta)$, the Laplace transform of an exponential distribution of pD

with rate β. So $D/t_{\mathrm{mem}} = 2pD\,[1 + O(p)]$ is exponential with mean s in the limit. Because both covariances are multiples of the identity to leading order, the same holds for the share of the capacity of any readout. Finally, $\sum_k kG_k = (1-\mathcal{F})^{-1}\Sigma = \Sigma/(\beta p) + O(p)$ gives $\mathbb{E}D = s\,t_{\mathrm{mem}} + O(1)$.

Two technical points. On symmetric matrices, where covariances live, the eigenvalue one of $\mathcal{F}_0$ is simple: a fixed point satisfies $\mathcal{T}_0 X = X$ and is therefore a multiple of the identity, and $\ell(1) > 0$ rules out a Jordan chain. For $p > 0$ every eigenvalue of $\mathcal{F}$ lies inside the unit circle, by a Perron–Frobenius argument: $\mathcal{F}$ maps positive matrices to positive matrices, and so does $(1 - \mathcal{F})^{-1} = 1 + (1 - \mathcal{T})^{-1}\mathcal{K}$. The eigenvalues of $\mathcal{F}_0$ therefore lie in the closed unit disk, only the simple one equals one, and only the eigenvalue of Eq. (B5) makes the sum diverge.

B.3 Physical channels without a reset

A physical channel does not reset to a prescribed state, but the proof used the reset only through three facts: one covariance direction is slow, it leaks at a small rate, and the mean step forgets everything else. The same three steps therefore prove the law for general channels.

Setting. One step is $r_t = A_u(\epsilon) r_{t-1} + b_u(\epsilon)$ on the coordinates that the inputs can reach, with a small leak ϵ, and $\mathcal{B}$, $\mathcal{T}$ and $\mathcal{K}$ are formed from A_u as before. We assume:

(i) A_u and b_u depend smoothly on ϵ in mean square, and b_u vanishes at $\epsilon = 0$.

(ii) The mean step forgets at $\epsilon = 0$: $\mathbb{E}A_u(0)$ has spectral radius below one.

(iii) The channel is mean-square stable for small $\epsilon > 0$: the covariance step $\mathcal{T}$ has spectral radius below one.

(iv) One covariance direction is slow: $\mathcal{T}$ has a simple eigenvalue $1 - \kappa\epsilon + O(\epsilon^2)$ with $\kappa > 0$ and no other eigenvalue near one. At $\epsilon = 0$ its right eigenvector is a positive matrix Q, and its left eigenvector is a weight ℓ that is nonnegative on positive matrices, with $\ell(Q) = 1$.

(v) The input writes along it: the change written by one input, of order ϵ, has covariance $\epsilon^2 V_* + O(\epsilon^3)$ with $\ell(V_*) > 0$.

The reset model is the case $\epsilon = p$, $\kappa = 2$, $Q = 1$ and $\ell = \mathrm{Tr}\,/n$. Condition (iv) concerns the covariance that the random steps carry, not the number of slow modes: at a fixed amplitude the inputs turn the 99 slow dark modes of the opposite-phase channel differently in each step and spread a stored signal over all of them faster than it leaks, so that one covariance direction stays slow. Without (iv) one observable shows several order scales: two reset sectors with mixing fractions 1/4 and 3/4, selected by a classical flag, give products of more than $x\,t_{\mathrm{mem}}$ inputs the share $w\,e^{-4x} + (1-w)\,e^{-4x/3}$ with $w = 0.82$, which exact moments reproduce.

Mixing law without a reset. Assume (i)–(v) and take a fixed readout $y = Mr$ whose projected slow covariance MQM^{T} has full rank. Then the mixing fraction $s = \ell(\mathcal{K}_0 Q)$ is positive and at most one, and ϵD becomes exponentially distributed with rate κ/s. The mean order is $s/(\kappa\epsilon) + O(1)$: the mixing fraction times the memory lifetime $1/(\kappa\epsilon)$.

Proof. Steps 1 and 2 carry over. The stationary mean is of order ϵ by (i) and (ii), and the slow eigenvalue piles the written covariance up along Q,

$$\Sigma = \epsilon\,\frac{\ell(V_*)}{\kappa}\,Q + O(\epsilon^2),$$

and every fixed order carries a covariance of order ϵ^2. Here the readout condition enters: the shares of the orders in the capacity of the readout follow from Σ only if MQM^{T} has full rank. For one observable a this requires $a^{\mathsf{T}} Q a > 0$. In step 3, $\mathcal{F}_0$ has the same fixed points as $\mathcal{T}_0$, so Q spans them, and $\ell \circ (1 - \mathcal{B}_0)$ is its left eigenvector, with value $\ell(\mathcal{K}_0 Q) = s$ on Q. Differentiating $(1 - \mathcal{B})\mathcal{F} = \mathcal{K}$ at $\epsilon = 0$ and applying it to Q gives $(1 - \mathcal{B}_0)\mathcal{F}_0' Q = \mathcal{T}_0' Q$, whose weight $\ell(\mathcal{T}_0' Q) = -\kappa$ is the rate of the slow eigenvalue of $\mathcal{T}$. The order eigenvalue is therefore $1 - (\kappa/s)\epsilon + O(\epsilon^2)$, and the Laplace transform follows as before. It remains to show that s is positive, which also makes the eigenvalue simple. Because $\mathcal{K}_0$ maps positive matrices to positive matrices and ℓ is nonnegative on them, $\ell(\mathcal{B}_0 Y) \le \ell(\mathcal{T}_0 Y) = \ell(Y)$ for every positive Y, so $0 \le \ell(\mathcal{B}_0^j \mathcal{K}_0 Q) \le s \le 1$ for every j. If s vanished, the series $\mathcal{F}_0 = \sum_j \mathcal{B}_0^j \mathcal{K}_0$ would give $\ell(\mathcal{F}_0 Q) = 0$, which contradicts $\mathcal{F}_0 Q = Q$ and $\ell(Q) = 1$.

For three spins the certificates of Appendix D

check these conditions, and for five spins we check the law itself.

Numerical check. We compared the law with exact moments for 24 reset reservoirs and readouts in dimensions two to four. The largest deviation from the law shrinks in proportion to p, from 0.08 at $p = 3\%$ to 0.003 at $p = 0.1\%$. For five spins the share of the capacity in products of more than K inputs approaches $e^{-K/(s\,t_{\mathrm{mem}})}$: its largest deviation, over K and the reservoirs, falls from 0.33 and 0.43 at $s\,t_{\mathrm{mem}} \approx 1$ to 0.07 and 0.14 at $s\,t_{\mathrm{mem}} \approx 20$, with the opposite-phase channel and with independent decay.

B.4 The scaling limit

Setting. The amplitude now falls with the leak, $\eta^2 = \nu\epsilon$ at fixed ν, and uniformly in the input

$$\begin{aligned} A_u &= O_0\big[1 + \eta\, a_1(u) + \eta^2 a_2(u) - \epsilon Y\big] + o(\epsilon), \\ b_u &= \epsilon\, O_0 c_0 + o(\epsilon), \end{aligned} \tag{B7}$$

with a rotation O_0, $\mathbb{E}a_1 = 0$, and a step that is a rotation without the leak, so that $a_1^{\mathsf{T}} = -a_1$ and $a_2 + a_2^{\mathsf{T}} = a_1^2$. Weak independent decay has this form with $\epsilon = \gamma$, and the strong-loss generator of Eq. (A2) has it on the dark coordinates with $\epsilon = 1/\gamma$: O_0 is the step of the Hamiltonian at the mean input, a_1 and a_2 are its response to the input field, and Y and c_0 the leak. Let Π_0 project onto the vectors that O_0 leaves fixed and Π onto the symmetric matrices that commute with O_0, the slow covariances. For d energy levels with distinct transition frequencies these are the covariances of the $d-1$ populations and one variance for each coherence, $d(d-1)$ numbers. On them, with $\bar{a}_2 = \mathbb{E}a_2$,

$$\begin{aligned} \mathsf{L}X &= \Pi(YX + XY^{\mathsf{T}}), \\ \mathsf{S}X &= -\Pi(\bar{a}_2 X + X\bar{a}_2^{\mathsf{T}}), \\ \mathsf{K}X &= \Pi\,\mathbb{E}[a_1 X a_1^{\mathsf{T}}] \end{aligned} \tag{B8}$$

are the leak, what the inputs take from a stored product on average, and what they pass on to the next order, with $\operatorname{Tr}\mathsf{S}X = \operatorname{Tr}\mathsf{K}X$.

Scaling limit, precise form. Let every slow covariance decay: the eigenvalues of $\mathsf{L}+\nu(\mathsf{S}-\mathsf{K})$ have positive real parts. As $\epsilon \to 0$ at fixed ν:

(a) The stationary mean converges to the vector m_0 fixed by O_0 with $\Pi_0(Y - \nu\bar{a}_2)m_0 = \Pi_0 c_0$.

(b) With $V_1 = \mathbb{E}[a_1 m_0 m_0^{\mathsf{T}} a_1^{\mathsf{T}}]$ and $\mathsf{F}_\nu = \nu(\mathsf{L}+\nu\mathsf{S})^{-1}\mathsf{K}$,

$$\begin{aligned} \Sigma \to \Sigma_\nu &= \nu\,[\mathsf{L} + \nu(\mathsf{S}-\mathsf{K})]^{-1}\Pi V_1, \\ G_k \to \mathsf{F}_\nu^{k-1}\,&\nu\,(\mathsf{L}+\nu\mathsf{S})^{-1}\Pi V_1. \end{aligned}$$

(c) F_ν has a spectral radius $\varrho < 1$. For a readout $y = Mr$ with $M\Sigma_\nu M^{\mathsf{T}}$ of full rank, the distribution of D converges to that of (b), with a finite mean and the share of $k = 1$ for single inputs.

(d) As $\nu \to 0$, Σ_ν and ϱ vanish in proportion to ν. As $\nu \to \infty$, if the identity is the only slow covariance that the inputs leave unchanged, Σ_ν vanishes as $1/\nu$ and $\varrho/(1-\varrho) = \nu\,\mathbb{E}\|a_1\|_F^2/(2\operatorname{Tr}Y) + O(1)$, the $s\,t_{\mathrm{mem}}$ of the reset model.

Proof. Over a lifetime O_0 turns many times, so only what turns with it accumulates, and on that the leak and the inputs act at the same order ϵ. This is the weak-coupling limit of open quantum systems [51], taken for the covariance over input histories. It rests on a lemma: let R be orthogonal, P the projector onto the vectors it leaves fixed and $M_\epsilon = 1 - R - \epsilon RN + o(\epsilon)$. If $\Gamma = -PNP$ is invertible on the range of P, then $\epsilon M_\epsilon^{-1} \to \Gamma^{-1}P$. Indeed $PR = RP = P$, so in the blocks of P and $1-P$ the first diagonal block of M_ϵ is $\epsilon\Gamma + o(\epsilon)$, the off-diagonal blocks are of order ϵ, and the second tends to $1-R$, which is invertible there.

With $R = O_0$ the lemma gives (a), since $1-\bar{A} = 1 - O_0 - \epsilon O_0(\nu\bar{a}_2 - Y) + o(\epsilon)$. On symmetric matrices $\mathcal{R}_0 X = O_0 X O_0^{\mathsf{T}}$ is orthogonal, and Π projects onto its fixed points. With primes for the maps of Eq. (B8) without Π, Eq. (A3) has $\mathcal{B} = \mathcal{R}_0[1-\epsilon(\mathsf{L}'+\nu\mathsf{S}')]$, $\mathcal{K} = \epsilon\nu\,\mathcal{R}_0\mathsf{K}'$ and, because one input writes $\eta\,O_0 a_1 m_0$, $V = \epsilon\nu\,\mathcal{R}_0 V_1$, each up to $o(\epsilon)$. The lemma gives $\epsilon(1-\mathcal{B})^{-1} \to (\mathsf{L}+\nu\mathsf{S})^{-1}\Pi$ and $\epsilon(1-\mathcal{T})^{-1} \to [\mathsf{L}+\nu(\mathsf{S}-\mathsf{K})]^{-1}\Pi$, hence (b), and the order transfer $(1-\mathcal{B})^{-1}\mathcal{K}$ converges to a map whose nonzero eigenvalues are those of F_ν. The orders then fall geometrically at a common rate for small ϵ, which gives (c) once $\varrho < 1$.

The semigroups generated by $-(\mathsf{L}+\nu\mathsf{S})$ and $-[\mathsf{L}+\nu(\mathsf{S}-\mathsf{K})]$ keep positive matrices positive, as projections with the positive map Π of $X \mapsto e^{tW}Xe^{tW^{\mathsf{T}}}$, $W = \nu\bar{a}_2 - Y$, without and with the

positive map $\nu\mathsf{K}'$ added to its generator, and the first is bounded by the second. The second decays by assumption, so both generators have positive inverses, and $(1-\mathsf{F}_\nu)^{-1} = 1+\nu[\mathsf{L}+\nu(\mathsf{S}-\mathsf{K})]^{-1}\mathsf{K}$ is positive, which gives $\varrho < 1$ as in Appendix B.2. On covariances of fixed vectors $\mathsf{L}+\nu\mathsf{S}$ acts as $X \mapsto W_0 X + X W_0^{\mathsf{T}}$ with $W_0 = \Pi_0(Y - \nu\bar{a}_2)\Pi_0$, so W_0 is invertible and m_0 exists.

For (d), (b) is linear in ν to first order. For large ν, $\mathsf{F}_\nu \to \mathsf{S}^{-1}\mathsf{K}$, which leaves the identity unchanged because $\mathsf{S}1 = \mathsf{K}1$, with the left eigenvector $X \mapsto \operatorname{Tr}\mathsf{S}X$. To first order in $1/\nu$ its eigenvalue moves to $1 - 2\operatorname{Tr} Y/(\nu\,\mathbb{E}\|a_1\|_F^2)$. The mean m_0 falls as $1/\nu$, V_1 as $1/\nu^2$, and Σ_ν, which piles up along the identity as in Eq. (B3), as $1/\nu$. In the reset model $Y = 1$ and $s = \eta^2\,\mathbb{E}\|a_1\|_F^2/n$ to leading order.

C One qubit in closed form

A qubit relaxes towards $+X$ at a rate ϵ, through the jump $\sqrt{\epsilon}\,|+x\rangle\langle -x|$, is turned about Z by $u\theta$ per step by the inputs $u = \pm 1$, and is read along Y. The averaged step keeps the axes X and Y apart and the input-dependent part swaps them, so Y carries only odd orders: order $2j+1$ has the share $(1-\chi)\chi^j$ of its capacity, where χ is the factor by which the recursion (A3) returns covariance from Y through X to Y, and $\mathbb{E}D = (1+\chi)/(1-\chi)$. At a fixed angle with $\sin\theta \neq 0$, $1-\chi = 3\epsilon/\sin^2\theta + O(\epsilon^2)$ as $\epsilon \to 0$: ϵD becomes exponentially distributed with rate $3/(2\sin^2\theta)$, and the share of single inputs vanishes. If instead $\theta^2/\epsilon \to \nu$, Appendix B.4 applies with $O_0 = 1$ and gives

$$\chi \to \frac{\nu^2}{(\nu+2)(\nu+1)}, \qquad \operatorname{Var} Y \to \frac{4\nu}{(\nu+2)(3\nu+2)}. \tag{C1}$$

Single past inputs keep the share $1-\chi$. The signal is largest at $\nu = 2/\sqrt{3}$, where single inputs carry 80% of the capacity, and vanishes as $4/(3\nu)$ for $\nu \gg 1$ and as ν for $\nu \ll 1$. Direct propagation of the density matrix reproduces these expressions.

D Three spins at strong loss

What is certified. Take three spins with a collective channel of opposite phases, Eq. (D1), four Hamiltonians H1–H4 with couplings drawn from fixed seeds, and every Hamiltonian whose three couplings lie in a small box around theirs. As γ grows:

- At a fixed input amplitude, the order of the memory of X_1, divided by γ, becomes exponentially distributed, and the signal in X_1 falls as $1/\gamma$ (boxes of half-width 1/32).
- With the amplitude scaled as $8/\sqrt{\gamma}$, single past inputs keep a finite share of the capacity of X_1, and the range of delays that holds 90% of their recall grows in proportion to γ (boxes of half-width 1/512).

Model and check. The three spins evolve under

$$H_\eta(u) = \sum_{i<j} J_{ij} X_i X_j + \tfrac{1}{2}\sum_i Z_i + \tfrac{3+\eta u}{4}\sum_i X_i \tag{D1}$$

with one jump $\sqrt{\gamma}\,(\sigma_1^- - \sigma_2^- - \sigma_3^-)$, steps of $\Delta t = 1/2$ and signs $u = \pm 1$, at $\eta = 1$ or $\eta = 8/\sqrt{\gamma}$. The jump leaves three dark states, with six slow covariances. At strong loss one step of the full master equation is a rotation of them up to corrections of order $1/\gamma$, so the mixing law without a reset (Appendix B.3) gives the first statement and the scaling limit (Appendix B.4) the second, here with explicit error bounds for the full master equation. Both end in finitely many strict inequalities, which the archive `spin_certificates.zip` in the repository checks in interval arithmetic at 256 bits (ball arithmetic of Arb, python-flint 0.9.0), and the note `three_spin_limits.pdf` writes out the derivation. An independent program builds the full finite-rate maps and agrees with the certified values to better than one part in a million. The certified rate of the exponential law lies between 61 and 133, and with the scaled amplitude at $\gamma = 2^{16}$ single inputs keep between 8% and 81% of the capacity of X_1.

Why coherence carries the memory. In the limit the mean m_0 consists of populations of the energy levels of the Hamiltonian restricted to the dark states, and a new input turns a population into a coherence between two levels: V_1 lies in the coherences. Dephasing in this energy basis after every step removes them, and the stored signal falls to order ϵ^2, the collapse as $1/\gamma^2$ of Sec. 3.3. Numerically the variance of X_1 stays

between 10^{-4} and 2×10^{-3} from $\gamma = 2^8$ to 2^{18} when the coherences are kept and falls as γ^{-2}, by six orders of magnitude, when they are erased.

E Numerical details and supporting checks

This appendix backs the main text, section by section.

E.1 Memory lifetime and dark states (Sec. 3.1)

Figures smooth recall curves for display only, by averaging each delay with its neighbours within a fifth of the delay, and no sum uses the smoothed values. The maps of Fig. 2a,b are computed at $\gamma = 0.1$, 0.25 and 2^{-1} to 2^6 and interpolated along the delay. Figure 2c stops the equal-phase channel and the two collective channels at $\gamma = 64$, where their recall at a fixed amplitude starts to reach past the 300 delays we sum. Summed over all delays with exact moments (Appendix A.4), their capacity keeps rising up to $\gamma = 256$.

Both tests of the dark states act on the 100 dark coordinates. For the dark-state model we write $\mathcal{P} = RR^\flat$ with $R^\flat R = 1$, where R embeds the dark coordinates α, evolve α under Eq. (A2) and record the same 45 features of the full state $(R - \mathcal{S}\mathcal{H}(u)R/\gamma)\,\alpha$. The readout is trained on them exactly as for the full reservoir, and the model has no fitted parameter (Fig. 3a). For the erasure test we replace the dark coordinates after every step by their time average, $r \mapsto r - RR^\flat(r - \bar{r})$. The control erases the dark subspace rotated by a random product of single-spin unitaries. Recall is measured at delays 1 to 100 with 800 training, 200 validation and 800 test inputs, at $\gamma = 16$ (Fig. 3b).

The decay out of the dark states falls exactly as $1/\gamma$: for the equal-phase channel $\gamma\bar{\kappa}$ approaches $\kappa_c = 2.945$ from $\gamma = 512$ on, so t_f grows in proportion to γ, while it levels off near 56 steps with the opposite-phase channel and near 110 with pairwise decay. The second-order calculation of Eq. (A2) gives the same κ_c at five spins and 1.4, 2.1 and 3.4 at three, four and six spins: at every size we computed, the decay falls as $1/\gamma$, with a prefactor that depends on the size.

Figure 6 repeats the curves of Fig. 2c at four and six spins, with 30 and 63 features, on eight and four reservoirs. The opposite-phase channel

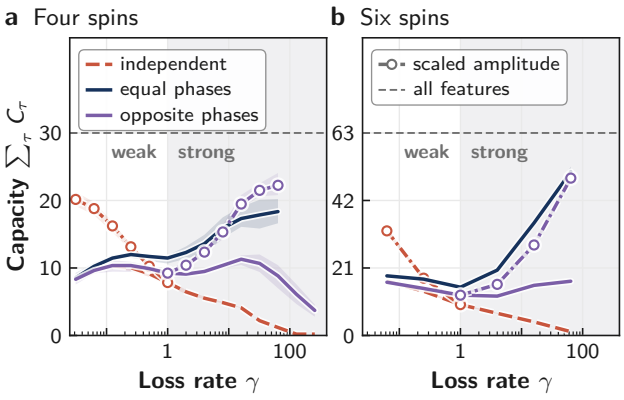


Figure 6: **Other reservoir sizes.** Capacity against the loss rate as in Fig. 2c, for four and six spins, with the same shading. A line stops where its recall reaches past the 300 delays summed.

reverses the sign of the last $\lfloor N/2 \rfloor$ spins. With the equal-phase channel the capacity grows with the loss at both sizes, from 11.5 to 18.4 units between $\gamma = 1$ and 64 at four spins and from 15 to 51 at six, and a scaled amplitude raises the best capacity of the other two channels. At three spins, which have only three dark states, the capacity of the equal-phase channel does not grow.

E.2 Recall time and phases (Secs. 3.2 and 3.3)

The recall time t_f shows the balance of decay and mixing. Its rate adds the decay of the stored signal and a mixing rate that grows as η^2 (Appendix A.3). The decay grows as γ with independent decay and falls as $1/\gamma$ with the opposite-phase channel, so $\eta^2 t_f$ depends on loss and input only through γ/η^2 and through $\gamma\eta^2$, the $1/\nu$ and ν of Appendix B.4. Both collapses hold for the four amplitudes $\eta = 1$, $\frac{1}{2}$, $\frac{1}{4}$ and $\frac{1}{8}$ (Fig. 7a), for γ from 2^{-8} to 2^6 with independent decay and from 2^{-2} to 2^{12} with the opposite-phase channel. With the slowest mode in place of the mean rate they hold to within 11% wherever the recall time exceeds four steps.

The write per input of Fig. 4b is $2\|R^\flat[E(\frac{3}{4}) - E(\frac{1}{4})]r_*\|$ at the stationary state r_* of the equal-phase channel, for the input patterns $(1, \dots, 1)$ and $\sqrt{2}\cos(2\pi i/5)$ multiplying X_i. The mode-matched pattern falls as $1/\gamma$, to 0.003 at $\gamma = 128$, while the other falls from 0.40 to 0.18. The rule is sharp: a phase error δ couples dark states directly at order δ, so beyond $\gamma \sim 1/\delta$ the equal-phase channel behaves like the opposite-phase channel.

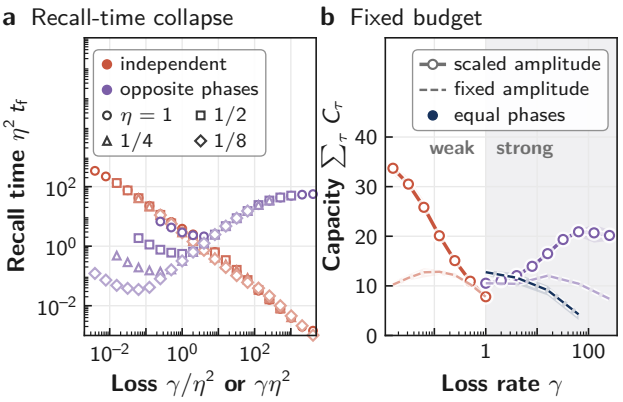


Figure 7: **Two supporting results.** **a** The time over which one input stays readable, rescaled as the sum of decay and mixing predicts: the four amplitudes fall on one curve per channel. **b** Capacity recovered at a budget of 10^9 reservoir steps per output step, with every replayed step counted. Shading as in Fig. 2c.

E.3 Inputs mixed and capacities (Secs. 3.2 and 3.3)

The distribution of the number of inputs mixed, Eq. (A3), uses six-point Gauss–Legendre quadrature for the expectations over the input and solves the equations $X - \bar{A}X\bar{A}^\mathsf{T} = C$ in the eigenbasis of $\bar{A}$. The shares of the 45 features are summed until the rest falls below 10^{-8}, or until successive orders shrink at a settled ratio, whose geometric tail we add. The linear recall follows from the same moments, $C_\tau = c_\tau^\mathsf{T}\Sigma_y^{-1}c_\tau/\mathrm{Var}\,u$ with c_τ the covariance of the features with $u_{t-\tau}$ and Σ_y their covariance, and reproduces the simulated capacities to a median of 2%. Three steps of inverse iteration on $\mathcal{T}$, from the input's covariance V and from the identity, give t_mem, Q and ℓ. Loss rates run in powers of four over the ranges of Fig. 4c, and with the scaled amplitude to 16384 and 1/4096, on four reservoirs, two at the longest memory lifetimes. At a fixed amplitude the mean order of the opposite-phase channel rises from 1.8 at $\gamma = 1$ to 23 at $\gamma = 4096$, and with independent decay it stays within 2.2 of $s\,t_\mathrm{mem}$ up to $s\,t_\mathrm{mem} = 18$.

Figure 4a shows the linear recall summed over all delays, from the same moments. At a fixed amplitude it peaks at 16.0 units with the opposite-phase channel ($\gamma = 64$, $s\,t_\mathrm{mem} = 0.32$) and at 14.9 with independent decay ($\gamma = 1/16$, 0.20), and falls to 2.2 and 3.7 units at lifetimes of 750 and 430 steps, where single past inputs carry 0.05 and 0.09 of the capacity. Table 3 sums the first 300 delays of simulated runs on eight reservoirs instead, which gives 14.7 units for the opposite-phase channel at $\gamma = 64$. With the scaled amplitude single inputs carry 0.82 of the capacity at a lifetime of 20 steps and 0.87 to 0.90 beyond 80 steps, the recall stays between 39 and 40 units, and $s\,t_\mathrm{mem}$ between 0.007 and 0.014. Of the inputs that the scaled amplitude removes from the mean of D with the opposite-phase channel, the mixing over the lifetime accounts for 43% at a lifetime of 20 steps, for 84% at 200 steps and for 99% at 750.

The scaling limit. We build the limit of Appendix B.4 from the generator of each reservoir: O_0, a_1 and a_2 are the orders zero to two of the step in the input field, and Y and c_0 its first order in the leak. Independent decay has 992 slow covariances and the opposite-phase channel 90, for the ten dark states under Eq. (A2), and at $\nu = 1$ they decay by at least $0.5\,\epsilon$ and $1.4\,\epsilon$ per step. With the opposite-phase channel 33 directions of the 45 features respond at leading order, 11 through the admixture of decaying operators of order $1/\gamma$, which the readout of the limit includes and which raises the share of single inputs by up to 0.02, and one at order $1/\gamma^2$, which the recursion no longer resolves from $\gamma = 1024$ on. The share of single inputs is here the share of $D = 1$, which includes the current input.

For the strong-loss generator the recursion converges onto the limit: at $\nu = 1$ the share of single inputs agrees to four digits from $\gamma = 2^{14}$ on. For the full master equation at $\gamma = 1/4096$ and 4096, on two reservoirs and for $\nu = 1/16$ to 64, it differs from the limit by at most 0.002 with independent decay and 0.009 with the opposite-phase channel, the mean of D by 0.003 and 0.016, and the signal, the variance that single inputs leave in the 45 features, by 0.4%. At $\nu = 1$ the limit gives single inputs 0.86 to 0.93 of the capacity over four reservoirs with independent decay and 0.90 to 0.94 with the opposite-phase channel. The signal is largest at $\nu = 22$ to 46 and at 68 to 215, where single inputs carry 0.38 to 0.55 of the capacity and $\varrho/(1-\varrho)$ lies between 0.8 and 4.8.

Table 3 collects the capacities. With exact readout all capacities of a readout add up to its rank, 45 here [29], so raising the linear memory capacity from 15 to 35 units leaves at most 10 instead of 30 for products of inputs.

Table 3: Linear memory capacity $\sum_\tau C_\tau$ of the five-spin reservoir in units (ceiling 45), summed over the first 300 delays of simulated runs: with exact readout, at a budget of 10^9 reservoir steps per output step, and with n shots per setting and step of replayed inputs (Sec. 3.4). Mean and standard deviation over eight reservoirs.

	independent decay	opposite phases	equal phases
tuned loss, $\eta = 1$[a]	14.5 ± 1.0	14.7 ± 1.1	30.6 ± 1.4
scaled amplitude[b]	34.9 ± 2.0	35.7 ± 3.0	–[c]
budget 10^9, $\eta = 1$[d]	12.8 ± 0.7	12.0 ± 0.7	12.8 ± 0.7
budget 10^9, scaled[e]	33.7 ± 1.1	20.9 ± 1.1	–[c]
$\gamma = 64$, exact[f]	1.2 ± 0.0	14.7 ± 1.1	30.6 ± 1.4
$\gamma = 64$, $n = 10^4$[f]	0.2 ± 0.0	8.2 ± 0.8	0.2 ± 0.0
$\gamma = 64$, $n = 10^6$[f]	0.2 ± 0.0	9.9 ± 0.7	2.3 ± 1.0

[a] At $\gamma = 0.1$, 64 and 64. [b] At $\gamma = 1/64$ and 256. [c] The uniform input is mode-matched to this channel. [d] At $\gamma = 1/8$, 16 and 1, the best loss rates on the grid. [e] At $\gamma = 1/64$ and 64. [f] At $\eta = 1$.

E.4 Best input amplitude (Secs. 3.3 and 3.4)

We mapped the linear memory capacity over loss rates and amplitudes: the first four reservoirs, independent decay at $\gamma = 2^{-8}$ to 2^{-2} and at 1, the opposite-phase channel at $\gamma = 1$, 4 and 2^3 to 2^8, and fifteen amplitudes from $\eta = 2^{-7}$ to 1, with exact features or 10^6 to 10^8 shots per setting and step from the state with the full history. With exact readout the capacity rises as the input weakens, up to 43 of the 45 units. With 10^6 to 10^8 shots an optimum exists. Its amplitude falls with the relaxation time, here the recall time at $\eta = 2^{-8}$, with fitted exponents between -0.40 and -0.46 in all six fits over the loss rates whose lifetime is at least four steps. Without the loss rates whose recall reaches past the 300 delays summed, the exponents are -0.48 to -0.56 with independent decay and -0.38 to -0.39 with the opposite-phase channel. The optimum is flat: one amplitude for all loss rates recovers on average 91% to 97% of the capacity at the optimum tracked at each loss rate, so the input need not be tuned to the loss rate.

E.5 Measurement (Sec. 3.4)

Weak signals cost measurements: with independent decay at $\gamma = 64$ the previous input moves the state by about 10^{-9}, which gives perfect recall of that input with exact readout and none with 10^8 shots per setting and step. With the opposite-phase channel, the scaled amplitude lowers the signal from 8×10^{-5} to 5×10^{-6}, and restarting each run from $|00000\rangle$, as the replay does, raises it to 4×10^{-5}, against 1.2×10^{-4} with independent decay and the scaled amplitude at $\gamma = 1/32$.

The replayed range τ_{90} grows with the memory lifetime: with the scaled amplitude of the opposite-phase channel from about 30 steps at $\gamma = 16$ to 130 at $\gamma = 256$. The thresholds of Sec. 3.4 are the crossings of the capacities at a fixed and with the scaled amplitude, from replayed runs at 10^3 to $10^{8.5}$ shots per setting and step. With replayed inputs the equal-phase channel at $\gamma = 64$ stays at the floor up to 10^4 shots and recovers 2.3 units at 10^6 and 15.5 at 3×10^8. There it passes the opposite-phase channel at a fixed amplitude, 11.0 units, and stays below it with the scaled amplitude, 26.3 units. At 10^3 and 10^4 shots the opposite-phase channel at $\gamma = 64$ recovers 5.4 and 8.2 units at a fixed amplitude, and independent decay at $\gamma = 1/32$ recovers 8.2 and 11.0. The scaled amplitude overtakes the fixed one at 2×10^4 shots with independent decay and at 10^5 with the opposite-phase channel, and independent decay then reaches 31.5 units at 3×10^8.

Figure 7b shows the capacity at a budget of 10^9 reservoir steps per output step, which is 3×10^6 to 3×10^7 shots per setting and step, and Table 3 lists its best values. With the full history, eight times more training inputs change the sampled capacity by less than 1.3 units.

For the matched relaxation time of Fig. 5b,c and Table 4 we find, for every channel and reservoir, the loss rate at which $t_{\mathrm{rel}} = 32$ steps, at weak loss when the kernel of the dissipator is the sta-

Table 4: The eight decay channels at a relaxation time of 32 input steps and $\eta = 0.35$: loss rate, stationary response, and linear memory capacity with exact readout and with 10^6 shots per measurement setting and step. Mean and standard deviation over eight reservoirs.

channel	loss rate γ	response	capacity exact	capacity at 10^6
independent	0.092	1.3	23.7 ± 1.1	17.4 ± 2.3
exchange	0.064	0.30	25.1 ± 0.8	10.1 ± 0.9
gain/loss	0.079	0.46	24.4 ± 0.8	12.5 ± 0.9
pairwise	168	1.5	18.2 ± 2.0	12.4 ± 1.2
opposite phases	101	0.68	25.3 ± 2.6	15.0 ± 1.5
two channels	73	0.011	26.6 ± 2.9	0.3 ± 0.1
equal phases	47	0.018	35.3 ± 0.8	0.3 ± 0.1
dephasing	143	0	0.0	0.2[a]

[a] A readout without any recall reads about 0.2 units, the floor of the sum over 300 delays.

tionary state alone and at strong loss otherwise, and run the reservoir there with $\eta = \sqrt{4/32}$. The stationary response is the Euclidean norm of the Pauli coefficients $\mathrm{Tr}(PW)$, over all Pauli strings P, of the change W: the commutator $-i[H_{\mathrm{in}}, \rho_*]$ with $H_{\mathrm{in}} = h\sum_i X_i$ and ρ_* the stationary state, projected spectrally onto the slow modes. These are the modes of the kernel, or without one those decaying at most four times faster than the slowest. The loss rate is matched for each reservoir, so Table 4 gives its mean. Over the 56 reservoirs of the seven channels that respond, the rank correlation of the stationary response is 0.96 with the recovered fraction of the capacity and 0.83 with the recovered capacity itself.

E.6 Ordinary noise

Adding independent decay $\sqrt{\gamma_{\mathrm{loc}}}\,\sigma_i^-$ to the equal-phase channel gives $\bar{\kappa} = \kappa_{\mathrm{c}}/\gamma + \kappa_{\mathrm{l}}\gamma_{\mathrm{loc}}$, where to first order κ_{l} is the trace of the added dissipator on the 99 slow dark modes divided by 99, $140/99 \approx 1.414$. Local dephasing $\sqrt{\gamma_\phi}\,Z_i$ adds $\kappa_\phi\gamma_\phi$ with $\kappa_\phi = 403.2/99 \approx 4.073$, and the full spectra reproduce both coefficients to within 0.2%. A recall time of at least T steps, $\kappa_{\mathrm{c}}/\gamma + \kappa_{\mathrm{l}}\gamma_{\mathrm{loc}} \le 1/(\Delta t\, T)$, then needs $\gamma/\gamma_{\mathrm{loc}} \ge 4\kappa_{\mathrm{c}}\kappa_{\mathrm{l}}\,\Delta t^2 T^2 \approx 4.2\,T^2$, or $12\,T^2$ for dephasing.

References

[1] Wolfgang Maass, Thomas Natschläger, and Henry Markram. "Real-time computing without stable states: A new framework for neural computation based on perturbations". Neural Computation **14**, 2531–2560 (2002).

[2] Herbert Jaeger. "The "echo state" approach to analysing and training recurrent neural networks". Technical Report 148. GMD Forschungszentrum Informationstechnik (2001).

[3] Keisuke Fujii and Kohei Nakajima. "Harnessing disordered-ensemble quantum dynamics for machine learning". Physical Review Applied **8**, 024030 (2017).

[4] Jiayin Chen and Hendra I. Nurdin. "Learning nonlinear input–output maps with dissipative quantum systems". Quantum Information Processing **18**, 198 (2019).

[5] Jiayin Chen, Hendra I. Nurdin, and Naoki Yamamoto. "Temporal information processing on noisy quantum computers". Physical Review Applied **14**, 024065 (2020).

[6] Antonio Sannia, Rodrigo Martínez-Peña, Miguel C. Soriano, Gian Luca Giorgi, and Roberta Zambrini. "Dissipation as a resource for quantum reservoir computing". Quantum **8**, 1291 (2024).

[7] Emanuele Ricci, Francesco Monzani, Luca Nigro, and Enrico Prati. "Quantum reservoir computing induced by controllable damping". npj Quantum Information **12**, 107 (2026).

[8] Laia Domingo, Gabriel Carlo, and Florentino Borondo. "Taking advantage of noise in quantum reservoir computing". Scientific Reports **13**, 8790 (2023).

[9] Niclas Götting, Steffen Wilksen, Alexander Steinhoff, Frederik Lohof, and Christopher Gies. "Connection between memory performance and optical absorption in quantum reservoir computing". Physical Review Letters **135**, 240403 (2025).

[10] Pere Mujal, Rodrigo Martínez-Peña, Gian Luca Giorgi, Miguel C. Soriano, and Roberta Zambrini. "Time-series quantum reservoir computing with weak and projective measurements". npj Quantum Information **9**, 16 (2023).

[11] Giacomo Franceschetto, Marcin Płodzień, Maciej Lewenstein, Antonio Acín, and Pere Mujal. "Harnessing quantum backaction for time-series processing". Physical Review X **16**, 021002 (2026).

[12] Kaito Kobayashi, Keisuke Fujii, and Naoki Yamamoto. "Feedback-driven quantum reservoir computing for time-series analysis". PRX Quantum **5**, 040325 (2024).

[13] Saud Čindrak, Brecht Donvil, Kathy Lüdge, and Lina Jaurigue. "Enhancing the performance of quantum reservoir computing and solving the time-complexity problem by artificial memory restriction". Physical Review Research **6**, 013051 (2024).

[14] Antonio Sannia, Ricard Ravell Rodríguez, Gian Luca Giorgi, and Roberta Zambrini. "Non-Markovianity and memory enhancement in quantum reservoir computing". npj Quantum Information **12**, 111 (2026).

[15] Iris Paparelle, Johan Henaff, Jorge García-Beni, Émilie Gillet, Daniel Montesinos, Gian Luca Giorgi, Miguel C. Soriano, Roberta Zambrini, and Valentina Parigi. "Experimental memory control in continuous-variable optical quantum reservoir computing". Nature Photonics **20**, 413–420 (2026).

[16] Rodrigo Martínez-Peña, Gian Luca Giorgi, Johannes Nokkala, Miguel C. Soriano, and Roberta Zambrini. "Dynamical phase transitions in quantum reservoir computing". Physical Review Letters **127**, 100502 (2021).

[17] Markus Baumann, Itamar Fink, Johannes Wittmann, Claudia Linnhoff-Popien, and Jonas Stein. "Where a quantum reservoir works: A transferable operating band" (2026). arXiv:2606.13284.

[18] J. F. Poyatos, J. I. Cirac, and P. Zoller. "Quantum reservoir engineering with laser cooled trapped ions". Physical Review Letters **77**, 4728–4731 (1996).

[19] Frank Verstraete, Michael M. Wolf, and J. Ignacio Cirac. "Quantum computation and quantum-state engineering driven by dissipation". Nature Physics **5**, 633–636 (2009).

[20] Fernando Pastawski, Lucas Clemente, and Juan Ignacio Cirac. "Quantum memories based on engineered dissipation". Physical Review A **83**, 012304 (2011).

[21] Patrick M. Harrington, Erich J. Mueller, and Kater W. Murch. "Engineered dissipation for quantum information science". Nature Reviews Physics **4**, 660–671 (2022).

[22] Robert H. Dicke. "Coherence in spontaneous radiation processes". Physical Review **93**, 99–110 (1954).

[23] A. Asenjo-Garcia, M. Moreno-Cardoner, A. Albrecht, H. J. Kimble, and D. E. Chang. "Exponential improvement in photon storage fidelities using subradiance and "selective radiance" in atomic arrays". Phys. Rev. X **7**, 031024 (2017).

[24] G. Facchinetti, S. D. Jenkins, and J. Ruostekoski. "Storing light with subradiant correlations in arrays of atoms". Phys. Rev. Lett. **117**, 243601 (2016).

[25] P. Facchi and S. Pascazio. "Quantum Zeno subspaces". Phys. Rev. Lett. **89**, 080401 (2002).

[26] Paolo Zanardi and Lorenzo Campos Venuti. "Coherent quantum dynamics in steady-state manifolds of strongly dissipative systems". Phys. Rev. Lett. **113**, 240406 (2014).

[27] Daniel Burgarth, Paolo Facchi, Hiromichi Nakazato, Saverio Pascazio, and Kazuya Yuasa. "Generalized adiabatic theorem and strong-coupling limits". Quantum **3**, 152 (2019).

[28] Florentin Reiter and Anders S. Sørensen. "Effective operator formalism for open quantum systems". Physical Review A **85**, 032111 (2012).

[29] Joni Dambre, David Verstraeten, Benjamin Schrauwen, and Serge Massar. "Information processing capacity of dynamical systems". Scientific Reports **2**, 514 (2012).

[30] Masanobu Inubushi and Kazuyuki Yoshimura. "Reservoir computing beyond memory-nonlinearity trade-off". Scientific Reports **7**, 10199 (2017).

[31] Saud Čindrak, Lara Giebeler, Niclas Götting, Christopher Gies, and Kathy Lüdge. "Memory-nonlinearity trade-off across quantum reservoir computing frameworks" (2026). arXiv:2603.21371.

[32] Fangjun Hu, Gerasimos Angelatos, Saeed A. Khan, Marti Vives, Esin Türeci, Leon Bello, Graham E. Rowlands, Guilhem J. Ribeill, and Hakan E. Türeci. "Tackling sampling noise in physical systems for machine learning applications: Fundamental limits and eigentasks". Physical Review X **13**, 041020 (2023).

[33] Olivia L. White, Daniel D. Lee, and Haim Sompolinsky. "Short-term memory in orthogonal neural networks". Phys. Rev. Lett. **92**, 148102 (2004).

[34] Surya Ganguli, Dongsung Huh, and Haim Sompolinsky. "Memory traces in dynamical systems". Proceedings of the National Academy of Sciences **105**, 18970–18975 (2008).

[35] Herbert Jaeger, Mantas Lukoševičius, Dan Popovici, and Udo Siewert. "Optimization and applications of echo state networks with leaky-integrator neurons". Neural Networks **20**, 335–352 (2007).

[36] Tomoyuki Kubota, Yudai Suzuki, Shumpei Kobayashi, Quoc Hoan Tran, Naoki Yamamoto, and Kohei Nakajima. "Temporal information processing induced by quantum noise". Physical Review Research **5**, 023057 (2023).

[37] Daniel Fry, Amol Deshmukh, Samuel Yen-Chi Chen, Vladimir Rastunkov, and Vanio Markov. "Optimizing quantum noise-induced reservoir computing for nonlinear and chaotic time series prediction". Scientific Reports **13**, 19326 (2023).

[38] Krai Cheamsawat and Thiparat Chotibut. "Dissipation alters modes of information encoding in small quantum reservoirs near criticality". Entropy **27**, 88 (2025).

[39] V. V. Akshay, M. V. Altaisky, and N. E. Kaputkina. "Quantum reservoir networks based on decoherence-free subspaces" (2026). arXiv:2605.27427.

[40] Vittorio Gorini, Andrzej Kossakowski, and E. C. George Sudarshan. "Completely positive dynamical semigroups of N-level systems". Journal of Mathematical Physics **17**, 821–825 (1976).

[41] Göran Lindblad. "On the generators of quantum dynamical semigroups". Communications in Mathematical Physics **48**, 119–130 (1976).

[42] Antonio Sannia, Gian Luca Giorgi, and Roberta Zambrini. "Exponential concentration and symmetries in quantum reservoir computing" (2025). arXiv:2505.10062.

[43] Almut Beige, Daniel Braun, Ben Tregenna, and Peter L. Knight. "Quantum computing using dissipation to remain in a decoherence-free subspace". Physical Review Letters **85**, 1762–1765 (2000).

[44] Maximilian Zanner, Tuure Orell, Christian M. F. Schneider, Romain Albert, Stefan Oleschko, Mathieu L. Juan, Matti Silveri, and Gerhard Kirchmair. "Coherent control of a multi-qubit dark state in waveguide quantum electrodynamics". Nature Physics **18**, 538–543 (2022).

[45] Mohammad Mirhosseini, Eunjong Kim, Xueyue Zhang, Alp Sipahigil, Paul B. Dieterle, Andrew J. Keller, Ana Asenjo-Garcia, Darrick E. Chang, and Oskar Painter. "Cavity quantum electrodynamics with atom-like mirrors". Nature **569**, 692–697 (2019).

[46] Rodrigo Martínez-Peña, Johannes Nokkala, Gian Luca Giorgi, Roberta Zambrini, and Miguel C. Soriano. "Information processing capacity of spin-based quantum reservoir computing systems". Cognitive Computation **15**, 1440–1451 (2023).

[47] Lyudmila Grigoryeva and Juan-Pablo Ortega. "Universal discrete-time reservoir computers with stochastic inputs and linear readouts using non-homogeneous state-affine systems". Journal of Machine Learning Research **19**, 1–40 (2018). url: https://jmlr.org/papers/v19/18-020.html.

[48] Ana Palacios, Rodrigo Martínez-Peña, Miguel C. Soriano, Gian Luca Giorgi, and Roberta Zambrini. "Role of coherence in many-body quantum reservoir computing". Communications Physics **7**, 369 (2024).

[49] Rodrigo Martínez-Peña and Juan-Pablo Ortega. "Quantum reservoir computing in finite dimensions". Physical Review E **107**, 035306 (2023).

[50] Rodrigo Martínez-Peña and Juan-Pablo Ortega. "Input-dependence in quantum reservoir computing". Physical Review E **111**, 065306 (2025).

[51] E. B. Davies. "Markovian master equations". Communications in Mathematical Physics **39**, 91–110 (1974).